

Big Data Technologies

SQL with Spark

Lionel Fillatre

Polytech Nice Sophia

lionel.fillatre@univ-cotedazur.fr


Outlines

- SparkSQL Concepts
- SparkSQL Guide
- Catalyst
- SparkSQL Conclusion

Spark SQL

Challenges and Solutions

- **Challenges**
 - Perform ETL (Extract-Transform-Load) to and from various (semi- or unstructured) data sources
 - Perform advanced analytics (e.g. machine learning, graph processing) that are hard to express in relational systems.
- **Solutions**
 - A DataFrame API that can perform relational operations on both external data sources and Spark's built-in RDDs.
 - A highly extensible optimizer, Catalyst, that uses features of Scala to add composable rule, control code generation, and define extensions.



About SQL

- Part of the core distribution since Spark 1.0 (April 2014)
- Spark Engine does not understand the structure of the data in RDDs or the semantics of user functions → limited optimization.
- Runs SQL / HiveQL queries, optionally alongside or replacing existing Hive deployments



```
SELECT COUNT(*)  
FROM hiveTable  
WHERE hive_udf(data)
```

Programming Interface




Figure 1: Interfaces to Spark SQL, and interaction with Spark.


SparkSQL Concepts

Various Data Sources Available in SparkSQL

- **Parquet Files** : It is a columnar format that is supported by many other data processing systems. Parquet files automatically preserves the schema of the original data.
- **ORC Files:** It is a free and open-source column-oriented data storage format (stores data tables by column rather than by row).
- **JSON Files:** It is an open-standard file format that uses human-readable text to transmit data objects consisting of attribute–value pairs and array data types (or any other serializable value)
- **Hive Tables**
- **JDBC To Other Databases**
- **Avro Files:** Avro is a data serialization system.


Apache Parquet File Format

- Apache Parquet is an open source file format that stores data in columnar format (as opposed to row format).
- Row-based formats such as CSV and JSON are (mostly) readable by humans, whereas column-based formats are optimized for computers.
- As a columnar file format, Apache Parquet can be read by computers much more efficiently and cost-effectively than other formats.




ORC File Format

- The Optimized Row Columnar (ORC) file format provides a highly efficient way to store Hive data.
- An ORC file contains groups of row data called **stripes**




What is Serialization?

- Serialization is the process of translating data structures or objects state into binary or textual form to transport the data over network or to store on some persistent storage.
- Once the data is transported over network or retrieved from the persistent storage, it needs to be deserialized again.
- Serialization is termed as **marshalling** and deserialization is termed as **unmarshalling**.



DataFrame

- A Dataset is a distributed collection of data
- A DataFrame is a Dataset organized into named columns.
- It is conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations under the hood. It is a distributed collection of rows with the same schema.
- Can be constructed from external data sources or RDDs into essentially an RDD of Row objects
- Supports relational operators (e.g. where, groupby) as well as Spark operations.
- Evaluated lazily ⇒ unmaterialized logical plan



Example of Dataframe

- **employee.json**

```
{  
  {"id": "1201", "name": "satish", "age": "25"},  
  {"id": "1202", "name": "krishna", "age": "28"},  
  {"id": "1203", "name": "amith", "age": "39"},  
  {"id": "1204", "name": "javed", "age": "23"},  
  {"id": "1205", "name": "prudvi", "age": "23"}  
}
```

- Code to create the dataframe

```
scala> val dfs = sqlContext.read.json("employee.json")  
dfs: org.apache.spark.sql.DataFrame = [age: string, id: string,  
name: string]
```

```
scala> dfs.show()  
+---+---+---+  
| age | id | name |  
+---+---+---+  
| 25 | 1201 | satish |  
| 28 | 1202 | krishna |  
| 39 | 1203 | amith |  
| 23 | 1204 | javed |  
| 23 | 1205 | prudvi |  
+---+---+---+
```

Data Model

- Spark SQL uses a nested data model based on Hive for tables and DataFrames
- Supports both primitive SQL types (boolean, integer, double, decimal, string, date, timestamp) and complex types (structs, arrays, maps, and unions); also user defined types.
- Complex data types can also be nested together to create more powerful types
- Accurately model data from a variety of sources and formats, including Hive, relational databases, JSON, and native objects in Java/Scala/Python

DataFrame Operations

- Relational operations (select, where, join, groupBy) via a domain-specific language (DSL) like SQL

employees

```
.join(dept, employees("deptId") === dept("id"))
.where(employees("gender") === "female")
.groupBy(dept("id"), dept("name"))
.agg(count("name"))
```

- Operators take *expression* objects
- Operators build up an Abstract Syntax Tree (AST), which is then optimized by *Catalyst*.
- Alternatively, register as temp SQL table and perform traditional SQL query strings

```
users.where(users("age") < 21)
.registerTempTable("young")
ctx.sql("SELECT count(*), avg(age) FROM young")
```

Advantages over Relational Query Languages

- DataFrames provide the same operations as relational query languages like SQL, we found that they can be significantly easier for users to work with thanks to their integration in a full programming language (control structures, e.g. if, for, etc.)
- Users can break up their code into Scala, Java or Python functions that pass DataFrames between them to build a logical plan
- Holistic optimization across functions composed in different languages.
- Logical plan analyzed eagerly ⇒ identify code errors associated with data schema issues on the fly.

Querying Native Datasets

- To interoperate with procedural Spark code, Spark SQL allows users to construct DataFrames directly against RDDs of objects native to the programming language
- Infer column names and types directly from data objects (via reflection in Java and Scala and data sampling in Python, which is dynamically typed)

```
case class User(name: String , age: Int)

// Create an RDD of User objects

usersRDD = spark.parallelize(
    List(User("Alice", 22), User("Bob", 19)))

// View the RDD as a DataFrame

usersDF = usersRDD.toDF
```

- Native objects accessed in-place to avoid expensive data format transformation.
- Benefits:
 - Run relational operations on existing Spark programs
 - Combine RDDs with external structured data

User-Defined Functions (UDFs)

- Easy extension of limited operations supported.
- Allows inline registration of UDFs.
- Can be defined on simple data types or entire tables.
- UDFs available to other interfaces (JDBC/ODBC for instance) after registration.
- Example:

```
val model: LogisticRegressionModel = ...
ctx.udf.register("predict",
  (x: Float, y: Float) => model.predict(Vector(x, y)))
ctx.sql("SELECT predict(age, weight) FROM users")
```

Spark SQL Guide

Datasets and DataFrames

- A Dataset is a distributed collection of data.
 - Dataset provides the benefits of RDDs (strong typing, ability to use powerful lambda functions) with the benefits of Spark SQL's optimized execution engine.
 - The Dataset API is available in Scala and Java. Python does not have the support for the Dataset API.
- A DataFrame is a Dataset organized into named columns.
 - It is conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations under the hood.
 - DataFrames can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing RDDs.
 - The DataFrame API is available in Scala, Java, Python, and R. In Scala and Java, a DataFrame is represented by a Dataset of Rows.
 - In the Scala API, DataFrame is simply a type alias of Dataset[Row].
- Scala/Java Datasets of Rows are often referred as DataFrames.

Creating DataFrames

- Create a dataframe from a json file
 - `val df = spark.read.json("/tmp/people.json")`
- Displays the content of the DataFrame to stdout
 - `df.show()`

```
// +---+-----+
// | age| name|
// +---+-----+
// |null|Michael|
// | 30| Andy|
// | 19| Justin|
// +---+-----+
```

people.json

```
{"name":"Michael"}
{"name":"Andy", "age":30}
 {"name":"Justin", "age":19}
```

DataFrame Operations

- This import is needed to use the \$-notation
 - `import spark.implicits._`
- Print the schema in a tree format
 - `df.printSchema()`

```
// root
// |-- age: long (nullable = true)
// |-- name: string (nullable = true)
```
- A **schema** is the description of the structure of your data (which together create a dataset in Spark SQL).
 - It can be **implicit** (and inferred at runtime) or **explicit** (and known at compile time).
 - The implicit schema can be exact or approximative.
 - A schema is described using StructType which is a collection of StructField objects (that in turn are tuples of names, types, and nullability classifier).

Untyped Dataset Operations

- Select only the "name" column

- df.select("name").show()

```
// +---+  
// | name|  
// +---+
```

```
// |Michael|
```

```
// | Andy|
```

```
// | Justin|
```

```
// +---+
```

- Select everybody, but increment the age by 1

- df.select(\$"name", \$"age" + 1).show()

```
// +---+-----+  
// | name|(age + 1)|  
// +---+-----+
```

```
// |Michael| null|
```

```
// | Andy| 31|
```

```
// | Justin| 20|
```

```
// +---+-----+
```

- Select people older than 21

- df.filter(\$"age" > 21).show()

```
// +---+---+  
// |age|name|  
// +---+---+
```

```
// | 30|Andy|  
// +---+---+
```

- Count people by age

- df.groupBy("age").count().show()

```
// +---+---+  
// | age|count|  
// +---+---+
```

```
// | 19| 1|  
// +---+---+
```

```
// |null| 1|  
// +---+---+
```

```
// | 30| 1|  
// +---+---+
```

Running SQL Queries Programmatically

- Register the DataFrame as a SQL temporary view

- df.createOrReplaceTempView("people")
- val sqlDF = spark.sql("SELECT * FROM people")
- sqlDF.show()

```
// +---+-----+
// | age| name|
// +---+-----+
// |null|Michael|
// | 30| Andy|
// | 19| Justin|
// +---+-----+
```

Global Temporary View

- Temporary views in Spark SQL are session-scoped and will disappear if the session that creates it terminates.
- If you want to have a temporary view that is shared among all sessions and keep alive until the Spark application terminates, you can create a global temporary view.
- Global temporary view is tied to a system preserved database `global_temp`, and we must use the qualified name to refer it, e.g. `SELECT * FROM global_temp.view1`.
- Register the DataFrame as a global temporary view
 - `df.createGlobalTempView("people")`
 - `spark.sql("SELECT * FROM global_temp.people").show()`

```
// +---+-----+
// | age|  name|
// +---+-----+
// |null|Michael|
// | 30|  Andy|
// | 19| Justin|
// +---+-----+
```

Creating Datasets

- Datasets are similar to RDDs, however, instead of using Java serialization or Kryo they use a specialized encoder to serialize the objects for processing or transmitting over the network.
- While both encoders and standard serialization are responsible for turning an object into bytes, encoders are code generated dynamically and use a format that allows Spark to perform many operations like filtering, sorting and hashing without deserializing the bytes back into an object.

- Define a case class

- `case class Person(name: String, age: Long)`

- Encoders are created for case classes

- `val caseClassDS = Seq(Person("Andy", 32)).toDS()`

- `caseClassDS.show()`

```
// +---+---+
```

```
// |name|age|
```

```
// +---+---+
```

```
// |Andy| 32|
```

```
// +---+---+
```

- The Case class in Scala is pretty much like a regular Scala class but with some additional functionality.
 - The objects of this class can be instantiated even without using the “new” keyword.
 - We can conveniently copy one object of the Case class to another entirely or even while changing some of the values of some of the attributes of this class.

Creating Datasets

- Encoders for most common types are automatically provided by importing `spark.implicits._`
 - `val primitiveDS = Seq(1, 2, 3).toDS()`
 - `primitiveDS.map(_ + 1).collect() // Returns: Array(2, 3, 4)`
- DataFrames can be converted to a Dataset by providing a class. Mapping will be done by name
 - `val path = "/tmp/people.json"`
 - `val peopleDS = spark.read.json(path).as[Person]`
 - `peopleDS.show()`

```
// +---+-----+
// | age|  name|
// +---+-----+
// |null|Michael|
// | 30|  Andy|
// | 19| Justin|
// +---+-----+
```

Programmatically Specifying the Schema


```
• import org.apache.spark.sql.Row // Apply the schema to the RDD
  • import org.apache.spark.sql.types._ // Create an RDD
    // Create an RDD
    • val peopleRDD = spark.sparkContext.textFile("/tmp/people.txt") // The schema is encoded in a string
      // The schema is encoded in a string
      • val schemaString = "name age" // Generate the schema based on the string of schema
        // Generate the schema based on the string of schema
        • val fields = schemaString.split(" ")
          .map(fieldName => StructField(fieldName, StringType, nullable = true))
        • val schema = StructType(fields) // Convert records of the RDD (people) to Rows
          // Convert records of the RDD (people) to Rows
          • val rowRDD = peopleRDD
            .map(_.split(","))
            .map(attributes => Row(attributes(0), attributes(1).trim()))
    • val peopleDF = spark.createDataFrame(rowRDD, schema) // Creates a temporary view using the DataFrame
      // Creates a temporary view using the DataFrame
      • peopleDF.createOrReplaceTempView("people") // SQL can be run over a temporary view created using DataFrames
        // SQL can be run over a temporary view created using DataFrames
        • val results = spark.sql("SELECT name FROM people") // The results of SQL queries are DataFrames and support all the normal RDD operations
          // The results of SQL queries are DataFrames and support all the normal RDD operations
          • results.map(attributes => "Name: " + attributes(0)).show()
            // +-----+
            // | value|
            // +-----+
            // |Name: Michael|
            // | Name: Andy|
            // | Name: Justin|
            // +-----+
    • results.map(attributes => "Name: " + attributes(0)).show()
```

Inferring the Schema from an RDD containing case classes

- The Scala interface for Spark SQL supports automatically converting an RDD containing case classes to a DataFrame.
- The case class defines the schema of the table.
- This RDD can be implicitly converted to a DataFrame and then be registered as a table.
- Tables can be used in subsequent SQL statements.
- For implicit conversions from RDDs to DataFrames
 - import spark.implicits._
- Create a class to model Person
 - `case class Person(name: String, age: Long)`
- Create an RDD of Person objects from a text file, convert it to a Dataframe
 - ```
val peopleDF = spark.sparkContext
 .textFile("examples/src/main/resources/people.txt")
 .map(_.split(","))
 .map(attributes => Person(attributes(0), attributes(1).trim.toInt))
 .toDF()
```
- Register the DataFrame as a temporary view
  - `peopleDF.createOrReplaceTempView("people")`
- SQL statements can be run by using the sql methods provided by Spark
  - `val teenagersDF = spark.sql("SELECT name, age FROM people  
 WHERE age BETWEEN 13 AND 19")`
- The columns of a row in the result can be accessed by field index
  - ```
teenagersDF.map(teenager => "Name: " + teenager(0)).show()  
// +-----+  
// | value |  
// +-----+  
// |Name: Justin|  
// +-----+
```

Catalyst


Plan Optimization & Execution



DataFrames and SQL share the same optimization/execution pipeline


Trees

- The main data type in Catalyst is a tree composed of node objects.
- Each node has a node type and zero or more children. New node types are defined in Scala as subclasses of the TreeNode class.
- These objects are immutable and can be manipulated using functional transformations.
- Example:
 - Literal(value: Int): a constant value
 - Attribute(name: String): an attribute from an input row, e.g., “x”
 - Add(left: TreeNode, right: TreeNode): sum of two expressions.
 - Add(Attribute(x), Add(Literal(1), Literal(2)))




Example of Tree

- Expression: `.where("tutorial='learntospark').limit(2)`




Catalyst Rules

- Pattern matching functions that transform subtrees into specific structures.
- Multiple patterns in the same transform call.
- May take multiple batches to reach a fixed point.
- Transform can contain arbitrary Scala code.



Analysis




Example:

```
SELECT col FROM sales
```

The type of *col*, or even whether it is a valid column name, is not known until we look up the table *sales*.

- Spark SQL begins with a relation to be computed, either from a tree returned by a SQL parser, or from a DataFrame object constructed using the API
- An attribute is **unresolved** if its type is not known or it's not matched to an input table.
- The Catalog object that tracks the tables in all data sources
- To resolve attributes:
 - Look up relations by name from the catalog.
 - Map named attributes to the input provided given operator's children.
 - Identifier for references to the same value
 - Propagate and coerce types through expressions

Logical Optimization




- Applies standard rule-based optimization (constant folding, predicate-pushdown, projection pruning, null propagation, boolean expression simplification, etc.)
- Example: when the fixed-precision DECIMAL type were added to Spark SQL, it was wanted to optimize aggregations such as sums and averages on DECIMALs with small precisions; it took just a few lines of code to write a rule that finds such decimals in SUM and AVG expressions, and casts them to unscaled 64-bit LONGs, does the aggregation on that, then converts the result back.

```
object DecimalAggregates extends Rule[LogicalPlan] {  
    /** Maximum number of decimal digits in a Long */  
    val MAX_LONG_DIGITS = 18  
  
    def apply(plan: LogicalPlan): LogicalPlan = {  
        plan transformAllExpressions {  
            case Sum(e @ DecimalType.Expression(prec, scale))  
                if prec + 10 <= MAX_LONG_DIGITS =>  
                    MakeDecimal(Sum(UnscaledValue(e)), prec + 10, scale)  
        }  
    }  
}
```

A simplified version of the code.

Physical Planning

```
def add_demographics(events):  
    u = sqlCtx.table("users")                      # Load partitioned Hive table  
    events.join(u, events.user_id == u.user_id)      # Join on user_id  
  
events = add_demographics(sqlCtx.load("/data/events", "parquet"))  
training_data = events.where(events.city == "Melbourne") # City is initially a field of "users"  
                    .select(events.timestamp).collect()
```



Physical Plan

- Physical plan is nothing but the conversion of the optimized logical plan into a physical plan that can be executed in the cluster.
- After the physical plans are generated, the cost is estimated recursively for the tree end-to-end.
- Finally, Spark uses the Cost Based Optimization to select the best suited physical plan to execute in cluster based on the provided data source.

Code Generation

- The final phase of query optimization involves generating Java bytecode to run on each machine.
- Catalyst relies on a special feature of the Scala language, quasiquotes, to make code generation simpler.
- Catalyst transforms a tree representing an expression in SQL to an Abstract Syntax Tree (AST) for Scala code to evaluate that expression, and then compile and run the generated code.
- The strings beginning with q are **quasiquotes**, meaning that although they look like strings, they are parsed by the Scala compiler at compile time.
- With code generation, we can write a function to translate a specific expression tree to a Scala AST as follows:

```
def compile(node: Node): AST = node match {  
    case Literal(value) => q"$value"  
    case Attribute(name) => q"row.get($name)"  
    case Add(left, right) =>  
        q"${compile(left)} + ${compile(right)}"  
}
```

SparkSQL Conclusion

SparkSQL Conclusion

- Let developers create and run Spark programs faster:
 - Write less code
 - Read less data
 - Let the optimizer do the hard work
- DataFrames and SQL provide a common way to access a variety of data sources
- Lots of extensions by using User Defined Functions