

# Data Science

# Reinforcement Learning


Lionel Fillatre  
Polytech Nice Sophia  
[lionel.fillatre@univ-cotedazur.fr](mailto:lionel.fillatre@univ-cotedazur.fr)

# Outline of this Lecture

- Introduction
  - State-Value Function
  - State-Action Value Function
  - Dynamic Programming
  - Conclusion
- 
- Appendix: value iteration algorithm convergence


# Introduction

# Branches of Machine Learning



# Characteristics of Reinforcement Learning

- What makes reinforcement learning different from other machine learning paradigms?
  - There is no supervisor, only a reward signal
  - Feedback is delayed, not instantaneous
  - Time really matters (sequential, non i.i.d data)
  - Agent's actions affect the subsequent data it receives



# Some examples

- Playing games: Atari

<https://www.youtube.com/watch?v=V1eYniJORnk>

- Playing games: Super Mario


[https://www.youtube.com/watch?v=wfL4L\\_I4U9A](https://www.youtube.com/watch?v=wfL4L_I4U9A)

- Making pancakes with a robot

[https://www.youtube.com/watch?v=W\\_gxLKSsSIE](https://www.youtube.com/watch?v=W_gxLKSsSIE)

# Reinforcement Learning

State:  $S_t$   $\xrightarrow{\quad}$   $S_{t+1} \sim \mathcal{P}(\quad | S_t, A_t)$   
Reward:  $R_t$   $\xrightarrow{\quad}$   $R_{t+1} \sim \mathcal{R}(\quad | S_t, A_t, S_{t+1})$



Trajectory:  $S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$

- Basic idea:
  - Receive feedback in the form of **rewards**
  - Agent's utility is defined by the reward function
  - Must (learn to) act so as to **maximize expected rewards**
  - All learning is based on observed samples of outcomes!


# Agent-Environment Interface

- The agent and environment interact at each of a sequence of discrete time steps,  $t = 0, 1, 2, 3, \dots$
- At each time step  $t$ , the agent receives some representation of the environment's state,  $S_t$ , and on that basis selects an action,  $A_t$
- One time step later, in part as a consequence of its action, the agent receives a numerical reward,  $R_{t+1}$ , and finds itself in a new state,  $S_{t+1}$
- The agent and environment together thereby give rise to a sequence or trajectory that begins like this:

$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$

# Markov Decision Process (MDP): States and Actions

- Let us describe the dynamics of the environment.
- The set of all possible actions is  $\mathcal{A}$
- The set of all possible states is  $\mathcal{S}$
- For a given state  $s$ ,  $\mathcal{A}(s)$  denotes the actions possible from this state  $s$  (some actions in  $\mathcal{A}$  may be impossible)
- The state is a description of the environment in sufficient detail to determine its evolution.
  - Think of Newtonian physics.
  - Markov assumption: the state at time  $t + 1$  depends directly on the state and action at time  $t$ , but not on past states and actions.
- To describe the dynamics, we need to specify the state-transition probabilities  $\mathbb{P}(S_{t+1}|S_t, A_t)$ .
- In this lecture, we assume the state  $S_t$  is fully observable at any time, a highly nontrivial assumption.



$$\begin{aligned}\mathcal{S} &= \{\text{Sleep, Run, Ice Cream}\} \\ \mathcal{A} &= \{\text{next}\} \\ &\text{(only one action)}\end{aligned}$$

# MDP: States and Actions



- Suppose you're playing Breakout. What should be the set of states and actions?
  - states = locations of the blocks, actions = paddle command (Right or Left)?
  - states = set of pixels, actions = trajectory of the paddle?
- In general, the right granularity of states and actions depends on what you're trying to achieve.

# Full Dynamics Model

- **Full transition probabilities:**

- It is given by the full model

$$p(s', r | s, a) = \mathbb{P}(S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a)$$

- This conditional probability gives the probability to achieve a state  $s'$  with a reward  $r$  from the state  $s$  when taking the action  $a$ .
  - Technically speaking, the reward  $R_{t+1} = r$  is a random value that is function of  $s'$ ,  $s$  and  $a$ , i.e. it is defined by the four-argument probability density function  $\mathbb{P}(R_{t+1} = r | S_{t+1} = s', S_t = s, A_t = a)$
  - All the presented results can be extended to this very general setting.

- **State-transition probabilities:**

- The state-transition probabilities  $p(s' | s, a) = \mathbb{P}(S_{t+1} = s' | S_t = s, A_t = a)$  are obtained from

$$p(s' | s, a) = \sum_{r \in \mathcal{R}} p(s', r | s, a)$$

where  $\mathcal{R}$  is the set of all possible rewards

# MDP: Expected Reward

- We can also compute the expected rewards for state–action pairs  $r(s, a)$ :

$$r(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} r \ p(r|s, a) = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r|s, a)$$

- Two-argument function
- Appropriate to model a reward just depending on state and action
- We can also compute the expected rewards for state–action–next-state triples  $r(s, a, s')$

$$r(s, a, s') = \mathbb{E}[R_t | S_t = s, A_t = a, S_{t+1} = s'] = \sum_{r \in \mathcal{R}} r \ p(r|s, a, s') = \sum_{r \in \mathcal{R}} r \frac{p(s', r|s, a)}{p(s'|s, a)}$$

- Three-argument function
- Appropriate to model a reward just depending on state, action and next state

# MDP: Policies

- The way the agent chooses the action in each step is called a policy.

- Two main types:


- Deterministic policy:  $A_t \sim \pi(S_t)$  for some function  $\pi: \mathcal{S} \rightarrow \mathcal{A}$
- Stochastic policy:  $A_t \sim \pi(\cdot | S_t)$  for some function  $\pi: \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A})$ . (Here,  $\mathcal{P}(\mathcal{A})$  is the set of distributions over actions.)

- With stochastic policies, the distribution over rollouts, or trajectories, factorizes:

$$p(s_1, a_1, \dots, s_T, a_T) = p(s_1)\pi(a_1|s_1)p(s_2|s_1, a_1)\pi(a_2|s_2) \cdots p(s_T|s_{T-1}, a_{T-1})\pi(a_T|s_T)$$

- Note:

- The fact that policies need consider only the current state is a powerful consequence of the Markov assumption and full observability.
- If the environment is partially observable, then the policy needs to depend on the history of observations.



# MDP: Reward

- **Episode:** A sequence of interactions based on which the reward will be judged at the end

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots, A_{T-1}, R_T$$

- The **return**  $G_t$  is the total discounted reward from time-step  $t$


$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+1+k}$$

- The discount  $\gamma \in [0,1]$  tells how much we care about reward now versus reward later.
  - $\gamma$  close to 0 leads to “myopic” evaluation
  - $\gamma$  close to 1 leads to “far-sighted” evaluation (the sum blows up for infinite length)
  - $\gamma = 1$  leads to undiscounted return


# Definition of the reward matters

- How might you define a reward function for an agent learning to play a video game?
  - **Score:** Change in score (why not current score?)
  - **Novelty:** Some measure of novelty (this is sufficient for most Atari games!)
  - **Demonstration:** Learning from demonstrations allows us to avoid specifying a reward directly and instead just learn to imitate how a human would complete the task.
  - **Feedback:** We can also incorporate human feedback by evaluating the quality of episodes or even sharing control with the agent in an interactive manner.
  - **Transfer learning:** It may be possible to use transfer learning to train on many similar games, and infer a “common sense” reward function for this game.

# Example: Gridworld



Actions (move on the grid): north, south, east, west



- The cells of the grid correspond to the states of the environment.
- Reward:
  - Actions that would take the agent off the grid leave its location unchanged, but also result in a reward of  $-1$ .
  - Other actions result in a reward of  $0$ , except :
    - From state A, all four actions yield a reward of  $+10$  and take the agent to A'
    - From state B, all actions yield a reward of  $+5$  and take the agent to B'.
- What's the strategy to achieve max reward?

# Planning and Policy

	<b>Prediction</b>	<b>Control</b>
<b>Model-based Planning</b>	<ul style="list-style-type: none"><li>• Dynamic programming: Policy Evaluation</li></ul>	<ul style="list-style-type: none"><li>• Dynamic programming: Policy Iteration</li><li>• Dynamic programming: Value Iteration</li></ul>
<b>Model-free reinforcement learning</b>	<ul style="list-style-type: none"><li>• Monte Carlo prediction</li><li>• Temporal Difference</li></ul>	<ul style="list-style-type: none"><li>• Monte Carlo control</li><li>• Sarsa</li><li>• Q learning</li></ul>

- **Planning:** a fully specified MDP is already known.
- **Reinforcement Learning:** agent interacts with an environment with unknown dynamics, i.e., the environment is a black box that takes in actions and outputs states and rewards.
- **Control in a model-free setting:** this is what most people mean when they use the term « Reinforcement learning ».

# State-Value Function

# Value Function

- Value function is a prediction of future reward
- Used to evaluate the goodness/badness of states and therefore to select between actions
- The **state-value function**  $v_\pi(s)$  of an MDP is the expected return starting from state  $s$ , and then following policy  $\pi$


$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \sum_{k=0}^{\infty} \gamma^k \mathbb{E}_\pi[R_{t+1+k} | S_t = s]$$

- $\mathbb{E}_\pi[X]$  denotes the expected value of a random variable  $X$  given that the agent follows policy  $\pi$
- Very important question: **How to calculate the value function?**

# Example: Maze

- Rewards:  $-1$  per time-step
- Actions: N, E, S, W
- States: Agent's location

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$



# How to calculate the Value Function?

- Note that

$$v_{\pi}(S_{t+1} = s') = \mathbb{E}_{\pi}[G_{t+1}|S_{t+1} = s'] = \mathbb{E}_{\pi}[G_t|S_t = s'] = v_{\pi}(s')$$

- Hence, the value function does not depend on time  $t$  (because of the MDP model)
- Note also that

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) = R_{t+1} + \gamma G_{t+1}$$

- Hence, we get
- Approach to calculate  $v_{\pi}(s)$ :
  - Compute  $\mathbb{E}_{\pi}[R_{t+1}|S_t = s]$
  - Compute  $\mathbb{E}_{\pi}[G_{t+1}|S_t = s]$
  - Combine both the results

# How to calculate $\mathbb{E}_\pi[R_{t+1}|S_t = s]$ ?

- With the **full four-argument** model:

$$\begin{aligned}\mathbb{E}_\pi[R_{t+1}|S_t = s] &= \sum_{r \in \mathcal{R}} r \ p(r|s) \\ &= \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} \sum_{a \in \mathcal{A}} p(s', r, a|s) \\ &= \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} \sum_{a \in \mathcal{A}} p(s', r|s, a) \mathbf{p}(a|s) \\ &= \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} \sum_{a \in \mathcal{A}} p(s', r|s, a) \mathbf{\pi}(a|s) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} r p(s', r|s, a)\end{aligned}$$

because  $p(s', r|s, a) = \frac{p(s', r, a|s)}{p(a|s)}$  and  $\mathbf{p}(a|s) = \mathbf{\pi}(a|s)$

# How to calculate $\mathbb{E}_\pi[R_{t+1}|S_t = s]$ ?

- With the **two-argument reward function**

$$\begin{aligned}\mathbb{E}_\pi[R_{t+1}|S_t = s] &= \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} r p(s', r|s, a) \\ &= \sum_{a \in \mathcal{A}} \pi(a|s) r(s, a)\end{aligned}$$

by definition of  $r(s, a)$

- For simplicity, we use this form  $\mathbb{E}_\pi[R_{t+1}|S_t = s] = \sum_{a \in \mathcal{A}} \pi(a|s) r(s, a)$  in the rest of the lecture (so we do not exploit the full MDP dynamics).
- This expression is generally sufficient in practice since our policy is just a function  $\pi(a|s)$  that does not depend on  $s'$

# Calculation of $\mathbb{E}_\pi[G_{t+1}|S_t = s]$ (slide 1/2)

- First let us prove that  $\mathbb{E}_\pi[G_{t+1}|S_t = s] = \mathbb{E}_\pi[v_\pi(S_{t+1})|S_t = s]$
- The conditional expectation has the following property:
  - For random variables  $X, Y, Z$ , we have

$$\mathbb{E}[\mathbb{E}[X|Z, Y]|Y] = \mathbb{E}[X|Y]$$

- It follows that

$$\begin{aligned}\mathbb{E}_\pi[G_{t+1}|S_t = s] &= \mathbb{E}_\pi[\mathbb{E}_\pi[G_{t+1}|S_{t+1}, S_t = s]|S_t = s] \\ &= \mathbb{E}_\pi[\mathbb{E}_\pi[G_{t+1}|S_{t+1}]|S_t = s] \\ &= \mathbb{E}_\pi[v_\pi(S_{t+1})|S_t = s]\end{aligned}$$

since  $G_{t+1}$  does not depends on  $S_t$  when  $S_{t+1}$  is known.

# Calculation of $\mathbb{E}_\pi[G_{t+1}|S_t = s]$ (slide 2/2)

- Next, let us calculate that  $\mathbb{E}_\pi[v_\pi(S_{t+1})|S_t = s]$


$$\begin{aligned}\mathbb{E}_\pi[v_\pi(S_{t+1})|S_t = s] &= \sum_{s' \in \mathcal{S}} v_\pi(s') p(s'|s) \\ &= \sum_{s' \in \mathcal{S}} v_\pi(s') \sum_{a \in \mathcal{A}} p(s', a|s) \\ &= \sum_{s' \in \mathcal{S}} \sum_{a \in \mathcal{A}} v_\pi(s') p(s'|s, a) \color{red}{p(a|s)} \\ &= \sum_{s' \in \mathcal{S}} \sum_{a \in \mathcal{A}} v_\pi(s') p(s'|s, a) \color{red}{\pi(a|s)}\end{aligned}$$

$$\mathbb{E}_\pi[v_\pi(S_{t+1})|S_t = s] = \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s' \in \mathcal{S}} p(s'|s, a) v_\pi(s')$$

# Value Function: Bellman equation

- The foundation of many RL algorithms is the fact that value functions satisfy a recursive relationship, called the **Bellman equation**:

$$\begin{aligned}v_{\pi}(s) &= \mathbb{E}_{\pi}[G_t | S_t = s] \\&= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\&= \sum_{a \in \mathcal{A}} \pi(a|s) r(s, a) + \gamma \sum_{a \in \mathcal{A}} \pi(a|s) \sum_{s' \in \mathcal{S}} p(s'|s, a) v_{\pi}(s') \\&= \sum_{a \in \mathcal{A}} \pi(a|s) \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_{\pi}(s') \right]\end{aligned}$$



$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_{\pi}(s') \right]$$

# Value Function: Bellman operator

- Viewing  $v_\pi$  as a vector  $V$  (where entries correspond to states), define the Bellman backup operator  $T_\pi$ :

$$T_\pi V(s) = \sum_{a \in A} \pi(a|s) \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) V(s') \right]$$

- The Bellman equation can be seen as a fixed point of the Bellman operator

$$T_\pi v_\pi = v_\pi$$

# Exact solution of the Bellman equations

- Let  $\mathcal{S} = \{s_1, s_2, \dots, s_q\}$

$$v_\pi = \begin{pmatrix} v_\pi(s_1) \\ \vdots \\ v_\pi(s_q) \end{pmatrix}, \quad \mathcal{R}_\pi = \begin{pmatrix} \sum_{a \in \mathcal{A}} \pi(a|s_1) r(s_1, a) \\ \vdots \\ \sum_{a \in \mathcal{A}} \pi(a|s_q) r(s_q, a) \end{pmatrix}, \quad \mathcal{P}_\pi = \begin{pmatrix} \sum_{a \in \mathcal{A}} \pi(a|s_1) p(s_1|s_1, a) & \dots & \sum_{a \in \mathcal{A}} \pi(a|s_1) p(s_q|s_1, a) \\ \vdots & \ddots & \vdots \\ \sum_{a \in \mathcal{A}} \pi(a|s_q) p(s_1|s_q, a) & \dots & \sum_{a \in \mathcal{A}} \pi(a|s_q) p(s_q|s_q, a) \end{pmatrix}$$

- The set of Bellman equations

$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_\pi(s') \right]$$

becomes



$$v_\pi = \mathcal{R}_\pi + \gamma \mathcal{P}_\pi v_\pi$$

- If the **environment model, the reward function and the policy are known** and if **the involved matrices are small enough**, we get

$$v_\pi = (I - \gamma \mathcal{P}_\pi)^{-1} \mathcal{R}_\pi \quad (\text{if } I - \gamma \mathcal{P}_\pi \text{ is invertible})$$

# Fixed point algorithm: Principle

- The Dottie number  $d$  is the only real solution in  $[-1,1]$  to the equation  $\cos x = x$



- We get  $\cos^n x = x$  for any integer  $n$
- We can show that  $\cos x$  is a contraction mapping: there exists  $0 < q < 1$  such that
$$|\cos x - \cos y| < q|x - y| \text{ for all } (x, y) \in [-1,1]^2$$
- Hence, the Dottie number  $d$  is given by  $\lim_{n \rightarrow \infty} x_n = d$  where

$$\begin{aligned}x_{n+1} &= \cos x_n \\x_0 &\in \mathbb{R}\end{aligned}$$

# An Approximation: Iterative Policy Evaluation

- Let's use dynamic programming to find  $v_\pi$ .
  - This is an **approximation** since this is an iterative algorithm ( $v_k$  replaced  $v_\pi$  in the Bellman equations)
  - But it is **easier** and **faster**
- **Iterative Policy Evaluation:**
  - Start from an initial vector  $v_1$  for  $v_\pi$
  - For each  $k = 1, 2, \dots$ ,
$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_k(s') \right]$$
- Observe:
  - A fixed point of this update is exactly a solution of the optimal Bellman equation (proof later).
  - We're treating  $v_\pi$  as a vector with  $|S|$  entries.
  - This type of algorithm is an instance of dynamic programming.

# Full Algorithm: Iterative Policy Evaluation

- This algorithm returns an estimate  $V$  of  $v_\pi$
- Initialize:
  - A small threshold  $\theta > 0$
  - $V(s)$  arbitrarily for all  $s \in \mathcal{S}$  and  $V(\text{terminal})$  to 0
- Loop:
  - $\Delta \leftarrow 0$
  - Loop for each  $s \in \mathcal{S}$  :
    - $v \leftarrow V(s)$
    - $V(s) \leftarrow \sum_{a \in A} \pi(a|s) [r(s, a) + \gamma \sum_{s'} p(s'|s, a) V(s')]$
    - $\Delta \leftarrow \max\{\Delta, |v - V(s)|\}$
  - until  $\Delta < \theta$

# State-Action Value Function

# State-Action Value Function

- A closely related but usefully different function is the state-action value function
- The **action-value function**  $q_\pi(s, a)$  is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a] = \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+1+k} | S_t = s, A_t = a \right]$$

- If you knew  $q_\pi$ , how would you obtain  $v_\pi$ ?


$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a)$$

- This requires knowing the policy, so in general it's rather easy to recover  $v_\pi$  from  $q_\pi$ .
- If you knew  $v_\pi$ , how would you obtain  $q_\pi$ ?


$$q_\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_\pi(s')$$

- This requires knowing the dynamics, so in general it's not easy to recover  $q_\pi$  from  $v_\pi$ .

# Illustration



Value Function



State-Value Function  
(just some state-action pairs are shown)

Proof of  $v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s)q_\pi(s, a)$

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$$

$$= \sum_{a \in \mathcal{A}} \pi(a|S_t = s) \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$

$$= \sum_{a \in \mathcal{A}} \pi(a|s)q_\pi(s, a)$$

Proof of  $q_\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) v_\pi(s')$

$$\begin{aligned}
q_\pi(s, a) &= \mathbb{E}_\pi[G_t | S_t = s, A_t = a] \\
&= \mathbb{E}_\pi[R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) | S_t = s, A_t = a] \\
&= \mathbb{E}_\pi[R_{t+1} | S_t = s, A_t = a] + \gamma \mathbb{E}_\pi[q_\pi(S_{t+1}, A_{t+1}) | S_t = s, A_t = a] \\
&= r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \sum_{a' \in \mathcal{A}} q_\pi(s', a') p(s', a' | s, a) \\
&= r(s, a) + \gamma \sum_{s' \in \mathcal{S}} \sum_{a' \in \mathcal{A}} q_\pi(s', a') \pi(a' | s') p(s' | s, a) \\
&= r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) \sum_{a' \in \mathcal{A}} \pi(a' | s') q_\pi(s', a') \\
&= r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s' | s, a) v_\pi(s')
\end{aligned}$$

since  $v_\pi(s') = \sum_{a' \in \mathcal{A}} \pi(a' | s') q_\pi(s', a')$

# State-Action Value Function: Bellman equation

- $q_\pi$  satisfies a Bellman equation very similar to  $v_\pi$  (proof is similar):

$$q_\pi(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \left[ \sum_{a' \in \mathcal{A}} \pi(a'|s') q_\pi(s', a') \right]$$

- Viewing  $q_\pi$  as a vector (where entries correspond to state-action pairs), define the Bellman operator  $B_\pi$ :

$$B_\pi q_\pi(s, a) = q_\pi(s, a)$$

- The Bellman equation can be seen as a fixed point of the Bellman operator

$$B_\pi q_\pi = q_\pi$$

# Value Functions

- Value functions measure the goodness of a particular state or state/ action pair:
  - How good is for the agent to be in a particular state or execute a particular action at a particular state, for a given policy.
- Optimal value functions measure the best possible goodness of states or state/action pairs under all possible policies.

	State values	Action values
Prediction	$v_\pi$	$q_\pi$
Control	$v^*$	$q^*$

# Dynamic Programming

# Greedy Policy

- Suppose you're in state  $s$ . You get to pick one action  $a$ , and then follow (fixed) policy  $\pi$  from then on. What do you pick?

$$a = \operatorname{argmax}_{a' \in \mathcal{A}} q_{\pi}(s, a')$$

- You get the **deterministic greedy policy**  $\pi(s) = \operatorname{argmax}_{a' \in \mathcal{A}} q_{\pi}(s, a')$  for state  $s$ , i.e.,  
 $\pi(a|s) = 1$  if  $a = \pi(s)$ , otherwise  $\pi(a|s) = 0$

# Greedy Policy

- Remember that  $v_\pi(s) = \sum_{a \in A} \pi(a|s)q_\pi(s, a)$ . For a deterministic policy  $\pi$ , it follows that

$$v_\pi(s) = q_\pi(s, \pi(s))$$

- If a deterministic policy  $\pi^*$  is optimal, then it must be the case that **for any state  $s$ :**

$$\pi^*(s) = \operatorname{argmax}_{a' \in \mathcal{A}} q_{\pi^*}(s, a'),$$

otherwise there exists  $(s, u)$  such as  $\pi^*(s) < q_{\pi^*}(s, u)$ , hence you could improve the policy  $\pi^*$  by changing  $\pi^*(s) = u$ .


# Optimal Policy

- The optimal greedy policy  $\pi^*(s) = \operatorname{argmax}_{a' \in \mathcal{A}} q_{\pi^*}(s, a')$  satisfies for any state  $s$ :

$$\pi^*(s) = \operatorname{argmax}_{a' \in \mathcal{A}} q_{\pi^*}(s, a') = \operatorname{argmax}_{a' \in \mathcal{A}} \left( r(s, a') + \gamma \sum_{s' \in S} p(s'|s, a') v_{\pi^*}(s') \right)$$

- It is denoted  $\pi^* = \text{greedy}(v_{\pi^*})$
- This leads to the **policy iteration algorithm**

# Principle of Policy Iteration



- Policy evaluation: estimate  $v_{\pi_k}$
- Policy improvement: generate  $\pi_{k+1}$  based on  $v_{\pi_k}$

# Policy Iteration for estimating $\pi^*$

- 1 - Initialize:
  - A small threshold  $\theta > 0$
  - $V(s)$  arbitrarily for all  $s \in \mathcal{S}$  and  $V(\text{terminal})$  to 0
  - $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$
- 2- Policy Evaluation
  - Loop:
    - $\Delta \leftarrow 0$
    - Loop for each  $s \in \mathcal{S}$  :
      - $v \leftarrow V(s)$
      - $V(s) \leftarrow \sum_{a \in A} \pi(a|s) [r(s, a) + \gamma \sum_{s'} p(s'|s, a) V(s')]$
      - $\Delta \leftarrow \max\{\Delta, |v - V(s)|\}$
  - until  $\Delta < \theta$
- 3- Policy Improvement
  - $policy\_stable \leftarrow true$
  - For each  $s \in \mathcal{S}$ :
    - $old\_action \leftarrow \pi(s)$
    - $\pi(s) \leftarrow \operatorname{argmax}_{a' \in \mathcal{A}} (r(s, a') + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a') v_{\pi}(s'))$
    - If  $old\_action \neq \pi(s)$ , then  $policy\_stable \leftarrow false$
  - If  $policy\_stable$  is true,
    - then stop and return  $V \approx v_{\pi^*}$  and  $\pi \approx \pi^*$
  - else
    - go to 2

# A faster solution?

- One drawback to policy iteration is that each of its iterations involves policy evaluation, which may itself be a protracted iterative computation requiring multiple sweeps through the state set.
- If policy evaluation is done iteratively, then convergence exactly to  $v_\pi$  occurs only in the limit. Must we wait for exact convergence, or can we stop short of that?
- One important special case is when policy evaluation is stopped after just one sweep (one update of each state).

# Principle of Value Iteration

- From the optimal policy, it follows that

$$v_{\pi^*}(s) = q_{\pi^*}(s, \pi^*(s)) = \max_{a' \in \mathcal{A}} \left( r(s, a') + \gamma \sum_{s' \in S} p(s'|s, a') v_{\pi^*}(s') \right)$$


- The principle of the so called **value iteration algorithm** is to use this equation instead to the two-step computation (evaluation + improvement) in the policy iteration algorithm

# Value Iteration for estimating $\pi^*$

- Initialize:
  - A small threshold  $\theta > 0$
  - $V(s)$  arbitrarily for all  $s \in \mathcal{S}$  and  $V(\text{terminal})$  to 0
  - $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$
- Loop:
  - $\Delta \leftarrow 0$
  - Loop for each  $s \in \mathcal{S}$  :
    - $v \leftarrow V(s)$
    - $V(s) \leftarrow \max_{a' \in \mathcal{A}} (r(s, a') + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a')V(s'))$
    - $\Delta \leftarrow \max\{\Delta, |v - V(s)|\}$
  - until  $\Delta < \theta$
- Output a deterministic policy  $\pi \approx \pi^*$  such that

$$\pi(s) = \operatorname{argmax}_{a' \in \mathcal{A}} \left( r(s, a') + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a')V(s') \right)$$

# Example: Gridworld




$v_*$  =

22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

- Where there are multiple arrows in a cell, all of the corresponding actions are optimal.

$\pi_*$  =



# Conclusion

# Conclusion

- Reinforcement learning
  - One of three machine learning paradigms, alongside supervised learning and unsupervised learning.
  - Set of applications still being expanded
- Many reinforcement learning algorithms utilize dynamic programming techniques to derive exact solutions but
  - The state and action spaces to be small enough;
  - The dynamics of the model must be known.
- Learning the dynamics
  - Monte Carlo methods (actual or simulated interaction with an environment)

# Appendix

# Proof of convergence of the value iteration algorithm


- The proof of convergence is based on the fixed point algorithm.
- The following slides detail this approach.

# Optimal State-Action Value Function

- Bellman equation for optimal policy  $\pi^*$ :

$$q_{\pi^*}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) q_{\pi^*}(s', \pi^*(s'))$$

$$q_{\pi^*}(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \max_{a' \in \mathcal{A}} q_{\pi^*}(s', a')$$



- Now  $q_{\pi^*} = q^*$  is the optimal state-action value function, and we can rewrite the optimal Bellman equation without mentioning  $\pi^*$ :

$$q^*(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \max_{a' \in \mathcal{A}} q^*(s', a') = B^* q^*(s, a)$$

- Turns out this is sufficient to characterize the optimal policy. So we simply need to solve the fixed point equation  $B^* q^* = q^*$ , and then we can choose

$$\pi^*(s) = \operatorname{argmax}_{a' \in \mathcal{A}} q^*(s, a')$$

# Bellman Fixed Points

- Some interesting problems could be reduced to finding fixed points of Bellman backup operators:
  - Evaluating a fixed policy  $\pi$

$$B_\pi q_\pi = q_\pi$$

- Finding the optimal policy

$$B^* q^* = q^*$$

- Idea: keep iterating the operator over and over again

$q \leftarrow B_\pi q$  (policy evaluation)

$q \leftarrow B^* q$  (finding the optimal policy)

- We're treating  $q_\pi$  or  $q^*$  as a vector with  $|\mathcal{S}||\mathcal{A}|$  entries.
- This type of algorithm is an instance of dynamic programming.

# Finding the Optimal State-Action Function

- Let's use dynamic programming to find  $q^*$ .

- Value Iteration:

- Start from an initial function  $q_1$  for  $q_\pi$
  - For each  $k = 1, 2, \dots,$

$$q_k(s, a) = r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \max_{a' \in \mathcal{A}} q_k(s', a')$$

- Observe: a fixed point of this update is exactly a solution of the optimal Bellman equation

# Proof of Fixed Point Convergence

- We get

$$\begin{aligned}|B^*q_1(s, a) - B^*q_2(s, a)| &= \left| \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \max_{a' \in \mathcal{A}} q_1(s', a') \right] - \left[ r(s, a) + \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \max_{a' \in \mathcal{A}} q_2(s', a') \right] \right| \\ &= \gamma \left| \left[ \sum_{s' \in \mathcal{S}} p(s'|s, a) \left[ \max_{a' \in \mathcal{A}} q_1(s', a') - \max_{a' \in \mathcal{A}} q_2(s', a') \right] \right] \right| \leq \gamma \sum_{s' \in \mathcal{S}} p(s'|s, a) \left| \max_{a' \in \mathcal{A}} q_1(s', a') - \max_{a' \in \mathcal{A}} q_2(s', a') \right| \\ &\leq \gamma \max_{s' \in \mathcal{S}, a' \in \mathcal{A}} |q_1(s', a') - q_2(s', a')| \sum_{s' \in \mathcal{S}} p(s'|s, a) = \gamma \max_{s' \in \mathcal{S}, a' \in \mathcal{A}} |q_1(s', a') - q_2(s', a')| \\ &= \gamma \|q_1 - q_2\|_\infty\end{aligned}$$

where  $\|q_1 - q_2\|_\infty$  denotes the  $L^\infty$  norm.

- This is true for any  $(s, a)$ , so

$$\|B^*q_1 - B^*q_2\|_\infty \leq \gamma \|q_1 - q_2\|_\infty$$

- $B^*$  is a contraction map.

# Proof of Fixed Point Convergence

- From

$$\|B^*q_1 - B^*q_2\|_\infty \leq \gamma \|q_1 - q_2\|_\infty$$

It follows that

$$\|(B^*)^{(k)}q_1 - (B^*)^{(k)}q_2\|_\infty \leq \gamma^k \|q_1 - q_2\|_\infty$$

where  $(B^*)^{(k)}$  denotes the operator  $B^*$  iterated  $k$  times:  $B^* \circ B^* \circ \dots \circ B^*$

- Let  $q^*$  be a fixed point of  $B^*$ . Then for any  $q$ ,

$$\|(B^*)^{(k)}q - q^*\|_\infty \leq \gamma^k \|q - q^*\|_\infty$$

- Hence, iterated application of  $B^*$ , starting from any  $q$ , converges exponentially to a unique fixed point.