

# Data Science

# Reinforcement Learning

Lionel Fillatre  
Polytech Nice Sophia  
[lionel.fillatre@univ-cotedazur.fr](mailto:lionel.fillatre@univ-cotedazur.fr)

# Outline of this Lecture

- Quick Recap
- Monte Carlo Learning
- Monte Carlo Control
- On-Policy
- Off-Policy
- Conclusion

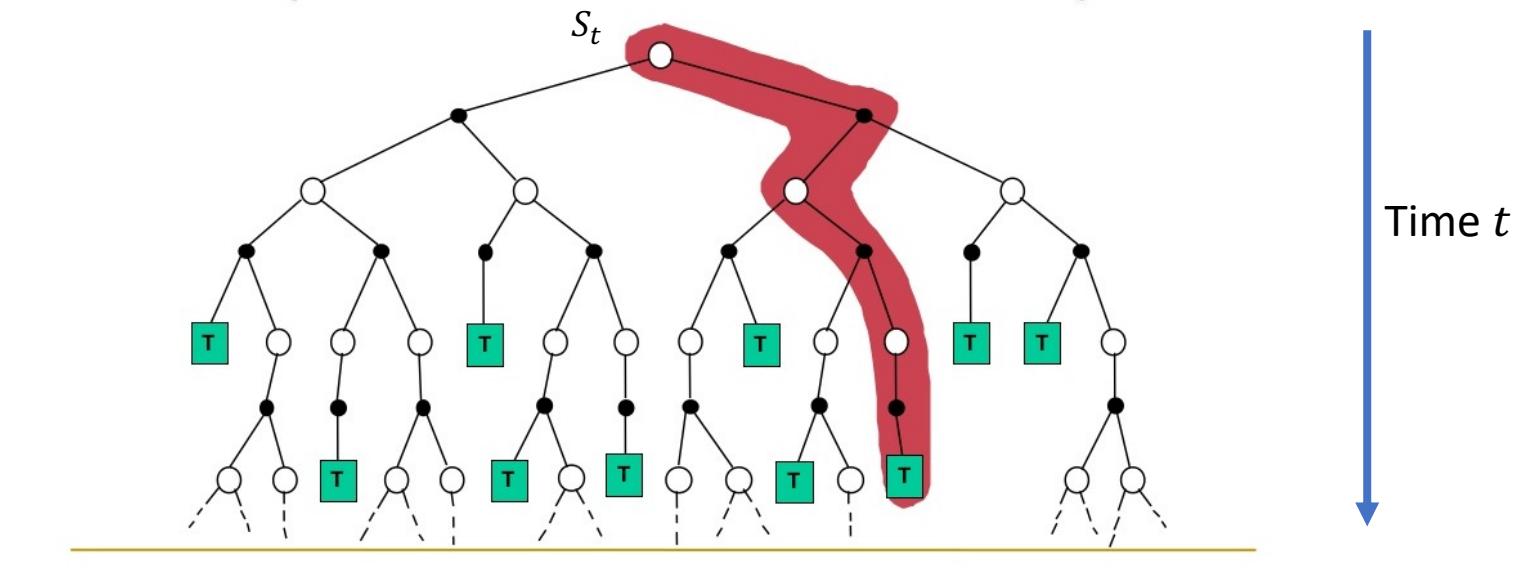
# Quick Recap: When dynamics are known...

- So far, we've focused on planning, where the dynamics are known.
- The optimal value functions are characterized in terms of a Bellman fixed point update.
- Since the Bellman operator is a contraction map, we can just keep applying it repeatedly, and we'll converge to a unique fixed point.
- **What are the limitations of value/policy iteration?**
  - Assumes known dynamics
  - The Markov property holds
  - Requires explicitly representing  $q^*$  or  $v^*$  as a vector
  - The number of states  $|\mathcal{S}|$  can be extremely large, or infinite
  - The number of actions  $|\mathcal{A}|$  can be infinite (e.g. continuous voltages in a machine)
- But value iteration is still a foundation for a lot of more practical RL algorithms.

# Monte Carlo Learning

# Monte Carlo (MC) Methods

- Monte Carlo methods are learning methods
  - Experience  $\Rightarrow$  values, policy
- Monte Carlo methods learn from complete sampled trajectories
  - Only defined for episodic tasks
  - All episodes must terminate
- Monte Carlo uses the simplest possible idea: value = mean return



# Monte-Carlo Policy Evaluation

- Remember that the value function is the expected return:

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

- Remember that the return is the total discounted reward:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots + \gamma^{T-1} R_T$$

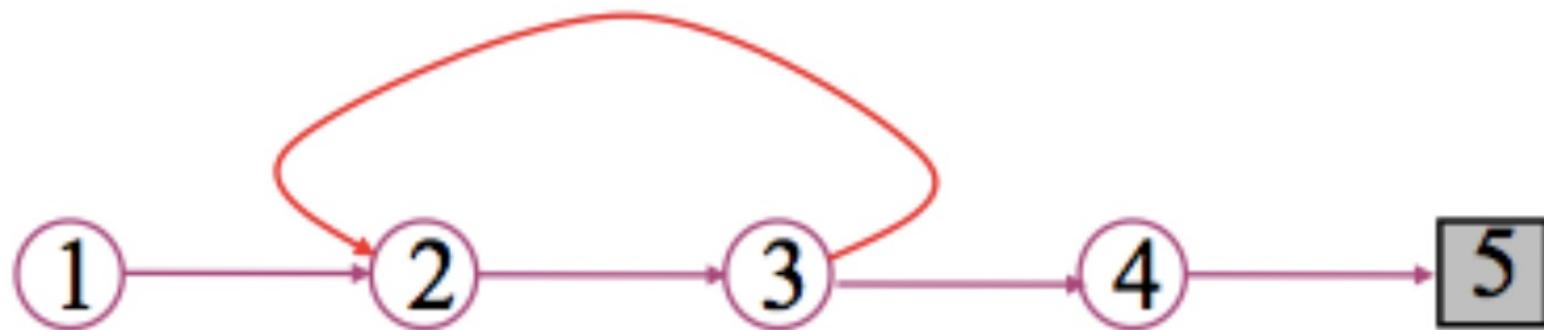
- Goal: learn  $v_{\pi}$  from episodes of experience under policy  $\pi$ :

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots, S_k \text{ when } A_1, A_2, A_3, \dots \sim \pi$$

- **Monte-Carlo policy evaluation uses empirical mean instead of expectation  $\mathbb{E}_{\pi}[\cdot]$**

# Monte-Carlo Policy Evaluation

- Goal: learn  $v_\pi(s)$  from episodes of experience under policy  $\pi$ :
- Idea: Average returns observed after visits to  $s$ :



- Every-Visit MC: average returns for every time  $s$  is visited in an episode
- First-visit MC: average returns only for first time  $s$  is visited in an episode
- Both converge asymptotically based on the law of large numbers

# Frozen Lake Environment

- The start state is  $s_1$ . The end state is the state  $s_{16}$ .
- The episode is described by the trajectory

$$S_0 = s_1, S_1 = s_5, S_2 = s_6, S_3 = s_7, S_4 = s_8, S_5 = s_{12}, S_6 = s_{11}, S_7 = s_{10}, S_8 = s_6, S_9 = s_7, S_{10} = s_8, S_{11} = s_{12}, S_{12} = s_{16}$$

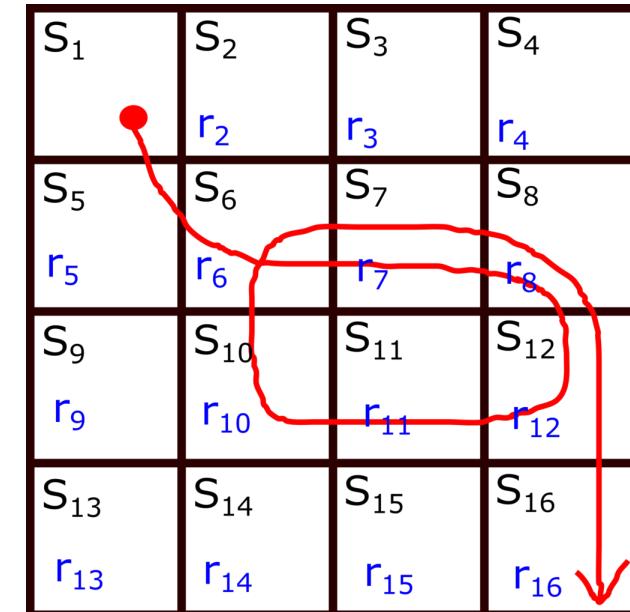
- Each occurrence of the state  $s_i$  in an episode is called a visit. There are two visits of  $s_7$
- The obtained rewards in the sequence are

$$R_1 = r_5, R_2 = r_6, R_3 = r_7, R_4 = r_8, R_5 = r_{12}, R_6 = r_{11}, R_7 = r_{10}, R_8 = r_6, R_9 = r_7, R_{10} = r_8, R_{11} = r_{12}, R_{12} = r_{16}$$

- Since we only have a single episode, the average return for every state will be equal to the return obtained at this single episode (first visit,  $\gamma = 1$ ).

$$\begin{aligned} \hat{v}_\pi(s_1) &= G(s_1)^{\text{episode 1}} = R_1 + R_2 + R_3 + R_4 + R_5 + R_6 + R_7 + R_8 + R_9 + R_{10} + R_{11} + R_{12} \\ &= r_5 + r_6 + r_7 + r_8 + r_{12} + r_{11} + r_{10} + r_6 + r_7 + r_8 + r_{12} + r_{16} \end{aligned}$$

$$\begin{aligned} \hat{v}_\pi(s_6) &= G(s_6)^{\text{episode 1}} = R_3 + R_4 + R_5 + R_6 + R_7 + R_8 + R_9 + R_{10} + R_{11} + R_{12} \\ &= r_7 + r_8 + r_{12} + r_{11} + r_{10} + r_6 + r_7 + r_8 + r_{12} + r_{16} \end{aligned}$$



Episode 1

$s_1$	$s_2$	$s_3$	$s_4$
$r_5$	$r_2$	$r_3$	$r_4$
$s_5$	$s_6$	$s_7$	$s_8$
$r_9$	$r_{10}$	$r_{11}$	$r_{12}$
$s_{13}$	$s_{14}$	$s_{15}$	$s_{16}$
$r_{13}$	$r_{14}$	$r_{15}$	$r_{16}$

Episode 2

$s_1$	$s_2$	$s_3$	$s_4$
$r_5$	$r_2$	$r_3$	$r_4$
$s_5$	$s_6$	$s_7$	$s_8$
$r_9$	$r_{10}$	$r_{11}$	$r_{12}$
$s_{13}$	$s_{14}$	$s_{15}$	$s_{16}$
$r_{13}$	$r_{14}$	$r_{15}$	$r_{16}$

Episode 3

$s_1$	$s_2$	$s_3$	$s_4$
$r_5$	$r_2$	$r_3$	$r_4$
$s_5$	$s_6$	$s_7$	$s_8$
$r_9$	$r_{10}$	$r_{11}$	$r_{12}$
$s_{13}$	$s_{14}$	$s_{15}$	$s_{16}$
$r_{13}$	$r_{14}$	$r_{15}$	$r_{16}$

- The returns from the state  $s_1$  for these three episodes is

$$G(s_1)^{\text{episode 1}} = r_5 + r_9 + r_{13} + r_{14} + r_{15} + r_{16}$$

$$G(s_1)^{\text{episode 2}} = r_5 + r_9 + r_{10} + r_{11} + r_{12} + r_{16}$$

$$G(s_1)^{\text{episode 3}} = r_5 + r_6 + r_7 + r_{11} + r_{15} + r_{16}$$

- Since the state  $s_1$  is visited 3 times, we have

$$\hat{v}_\pi(s_1) = \frac{G(s_1)^{\text{episode 1}} + G(s_1)^{\text{episode 2}} + G(s_1)^{\text{episode 3}}}{3}$$

- The returns from the state  $s_9$  for these three episodes is

$$G(s_9)^{\text{episode 1}} = r_{13} + r_{14} + r_{15} + r_{16}$$

$$G(s_9)^{\text{episode 2}} = r_{10} + r_{11} + r_{12} + r_{16}$$

No return for episode 3

- Since the state  $s_9$  is visited 2 times, we have

$$\hat{v}_\pi(s_9) = \frac{G(s_9)^{\text{episode 1}} + G(s_9)^{\text{episode 2}}}{2}$$

# First-Visit Monte Carlo Evaluation

- Input: a policy  $\pi$  to be evaluated
- Initialize:
  - $V(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$
  - $Returns(s) \leftarrow$  empty list, for all  $s \in \mathcal{S}$
- Loop forever (for each episode):
  - Generate an episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$
  - $G \leftarrow 0$
  - Loop for each step of episode,  $t = T - 1, T - 2, \dots, 0$ :
    - $G \leftarrow \gamma G + R_{t+1}$
    - If  $S_t$  does not appear in  $S_0, S_1, \dots, S_{t-1}$  (first visit of  $S_t$ ):
      - Append  $G$  to  $Returns(S_t)$
      - $V(S_t) \leftarrow \text{average}(Returns(S_t))$

# First-Visit Monte Carlo Evaluation

- To evaluate state  $s$  from several episodes
- The first time-step  $t$  that state  $s$  is visited in an episode
  - Incrementer counter for each episode:  $N(s) \leftarrow N(s) + 1$
  - Increment total return:  $G(s) \leftarrow G(s) + G_t$
- Value is estimated by mean return:  $\hat{v}(s) = \frac{G(s)}{N(s)}$
- By law of large numbers:  $\hat{v}(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

# Incremental Mean

- The mean  $\mu_k$  of a sequence  $x_1, \dots, x_k$  can be computed incrementally:

$$\begin{aligned}\mu_k &= \frac{1}{k} \sum_{j=1}^k x_j \\ &= \frac{1}{k} \left( x_k + \sum_{j=1}^{k-1} x_j \right) \\ &= \frac{1}{k} (x_k + (k - 1)\mu_{k-1}) \\ &= \mu_{k-1} + \frac{1}{k} (x_k - \mu_{k-1})\end{aligned}$$

# Incremental Monte Carlo Evaluation

- Update  $\hat{v}(s)$  incrementally after episode  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_k$  and visits
- For each state  $S_t$  with return  $G_t$ 
  - $N(S_t) \leftarrow N(S_t) + 1$
  - $\hat{v}(S_t) = \frac{G(S_t)}{N(S_t)} \leftarrow \hat{v}(S_t) + \frac{1}{N(S_t)}(G_t - \hat{v}(S_t))$
- In non-stationary problems, it can be useful to track a running mean, i.e. forget old episodes
  - $$\hat{v}(S_t) \leftarrow \hat{v}(S_t) + \alpha(G_t - \hat{v}(S_t))$$

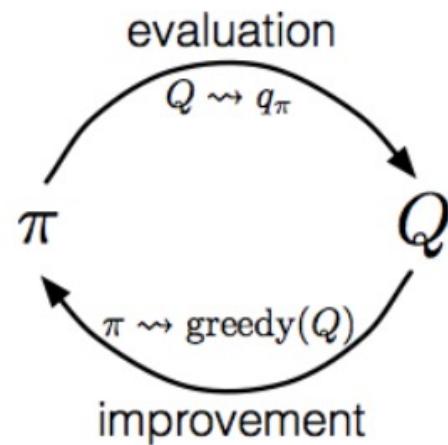
# Every-VISIT Monte Carlo Evaluation

- To evaluate state  $s$  from several episodes and from several visits
- Every time-step  $t$  that state  $s$  is visited in an episode
  - Incrementer counter for each episode:  $N(s) \leftarrow N(s) + 1$
  - Increment total return:  $G(s) \leftarrow G(s) + G_t$
- Value is estimated by mean return:  $\hat{v}(s) = \frac{G(s)}{N(s)}$
- By law of large numbers:  $\hat{v}(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

# Monte Carlo Control

# Monte-Carlo Control

$$\pi_0 \xrightarrow{E} q_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} q_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} q_*$$



Greedy policy:  $\pi(s) = \operatorname{argmax}_{a' \in A} Q(s, a')$

- **MC policy iteration step:** Policy evaluation using MC methods followed by policy improvement
- **Policy improvement step:** greedify with respect to value (or state-action value) function

# MC Estimation of Action Values

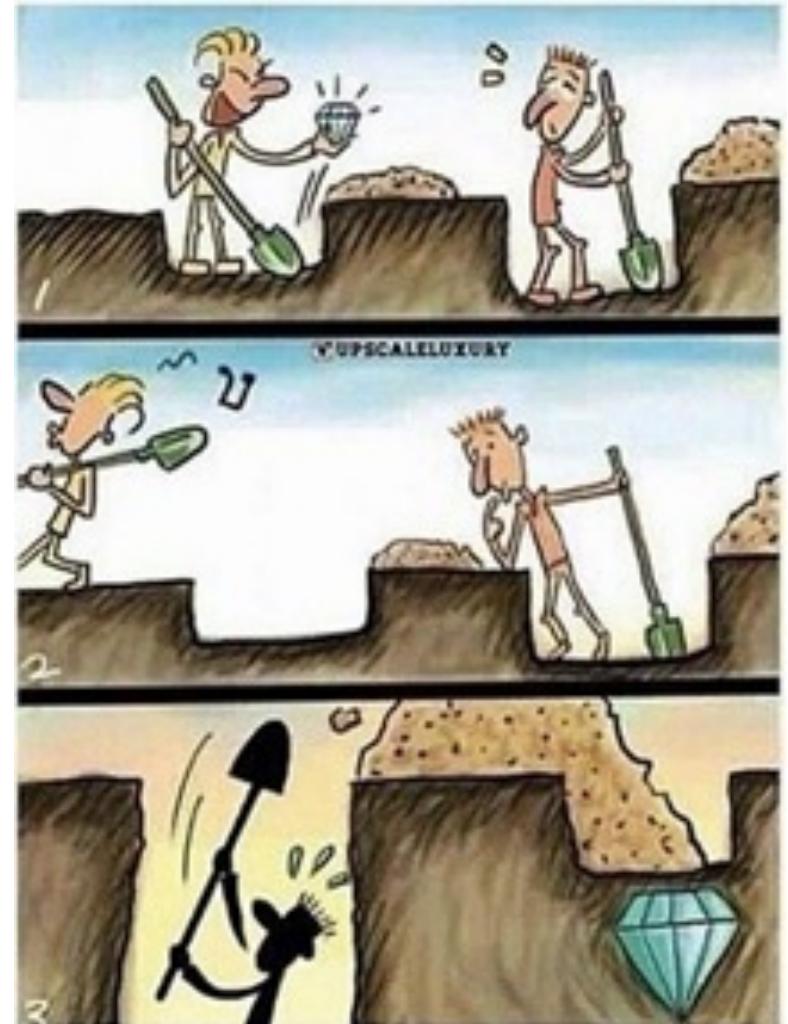
- MC is most useful when a model is not available
  - We want to learn  $q^*(s, a)$  because we can get an optimal policy without knowing dynamics
  - **Reminder:**  $q_\pi(s, a)$  - average return starting from state  $s$  and action  $a$  following  $\pi$

$$q_\pi(s, a) = r(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) v_\pi(s')$$

- Converges asymptotically if every state-action pair is visited.
- Is this possible if we are using a deterministic policy?

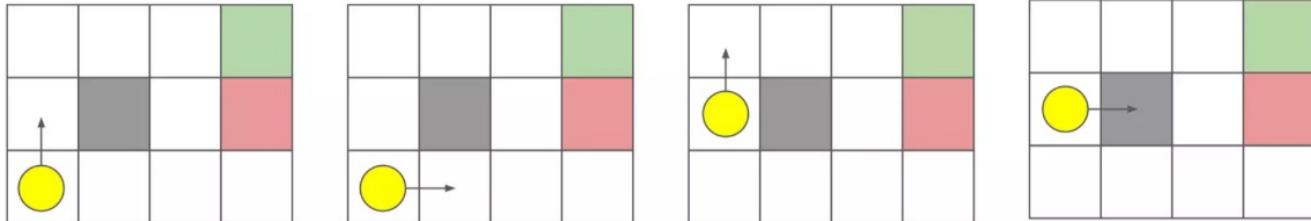
# The Exploration Problem

- If we always follow the deterministic policy to collect experience, we will never have the opportunity to see and evaluate (estimate  $q$ ) of alternative actions.
- All learning methods face a dilemma: they seek to learn action values conditioned on subsequent optimal behaviour but they need to act suboptimally in order to explore all actions (to discover the optimal actions).  
**This is the exploration-exploitation dilemma.**
- Solutions (discussed later):
  - **Exploring starts:** Every state-action pair has a non-zero probability of being the starting pair
  - Give up on deterministic policies and only search over  $\epsilon$ -soft policies
  - **Off-policy:** use a different policy to collect experience than the one you care to evaluate



# Monte Carlo Exploring Starts

- Initialize:
  - $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$
  - $Q(s, a) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$
  - $Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$



Several  $(S_0, A_0)$

- Loop forever (for each episode):
  - **Choose  $S_0 \in \mathcal{S}, A_0 \in \mathcal{A}(S_0)$  randomly such that all pairs have probability  $> 0$**
  - Generate an episode from  $S_0, A_0$ , following  $\pi$ :
$$S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$$
  - $G \leftarrow 0$
  - Loop for each step of episode,  $t = T - 1, T - 2, \dots, 0$ :
    - $G \leftarrow \gamma G + R_{t+1}$   
If the pair  $S_t, A_t$  does appear in  $S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}$ :
      - Append  $G$  to  $Returns(S_t, A_t)$
      - $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$
      - $\pi(S_t) \leftarrow \underset{u \in A}{\text{argmax}} Q(S_t, u)$

# Convergence of MC Control

- Greedified policy meets the conditions for policy improvement:

$$\begin{aligned} q_{\pi_k}(s, \pi_{k+1}(s)) &= q_{\pi_k}\left(s, \operatorname{argmax}_{a' \in \mathcal{A}} q_{\pi_k}(s, a')\right) \\ &= \max_{a' \in \mathcal{A}} q_{\pi_k}(s, a') \\ &\geq q_{\pi_k}(s, \pi_k(s)) \\ &= v_{\pi_k}(s) \end{aligned}$$

- And thus  $\pi_{k+1}$  must be better than  $\pi_k$  or as good as  $\pi_k$ .
- This assumes **exploring starts** and **infinite number of episodes** for MC policy evaluation

# Monte Carlo Control without Exploring Starts

- How can we avoid the unlikely assumption of exploring starts?
- The only general way to ensure that all actions are selected infinitely often is for the agent to continue to select them.
- There are two approaches to ensuring this, resulting in what we call **on-policy methods** and **off-policy methods**.
  - On-policy methods attempt to evaluate or improve the policy that is used to make decisions,
  - Off-policy methods evaluate or improve a policy different from that used to generate the data.

# On-Policy

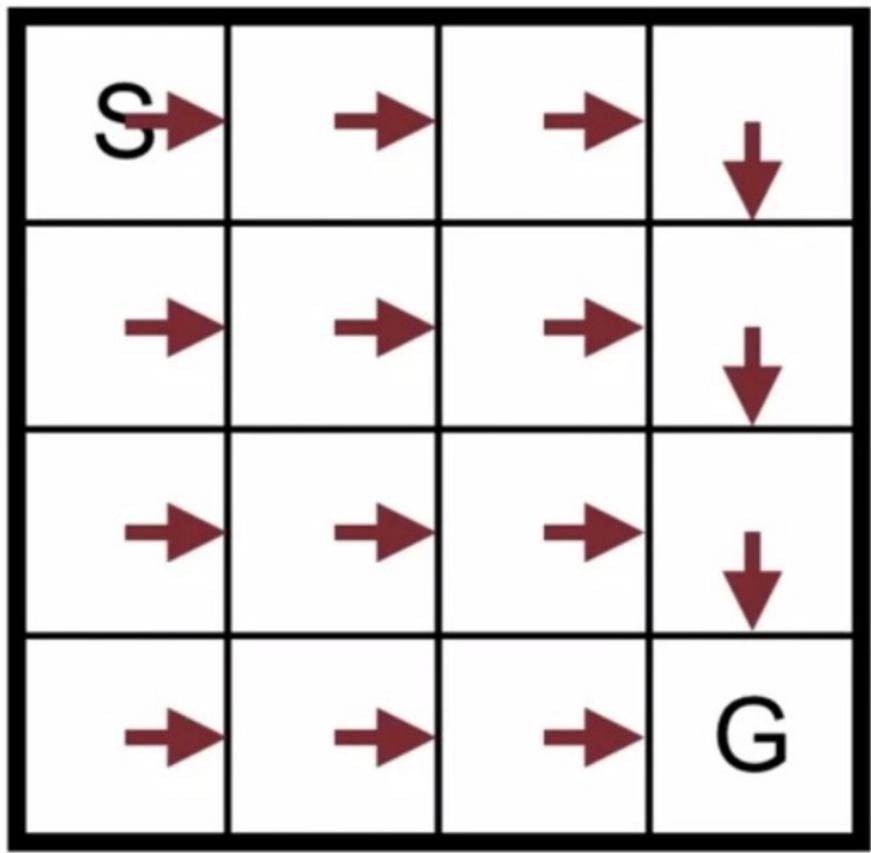
# On-policy Monte Carlo Control

- On-policy: learn about policy currently executing
- The policy must be eternally soft:  $\pi(a|s) > 0$  for all  $s$  and  $a \in \mathcal{A}(s)$  where  $\mathcal{A}(s)$  is the set of actions possible from state  $s$ .
- An  **$\epsilon$ -soft policy** is defined as a soft policy such as  $\pi(a|s) \geq \frac{\epsilon}{|\mathcal{A}(s)|}$  for all states and actions
- For example, let us consider the  **$\epsilon$ -greedy policy**. Its probability  $\pi(a|s)$  is defined by

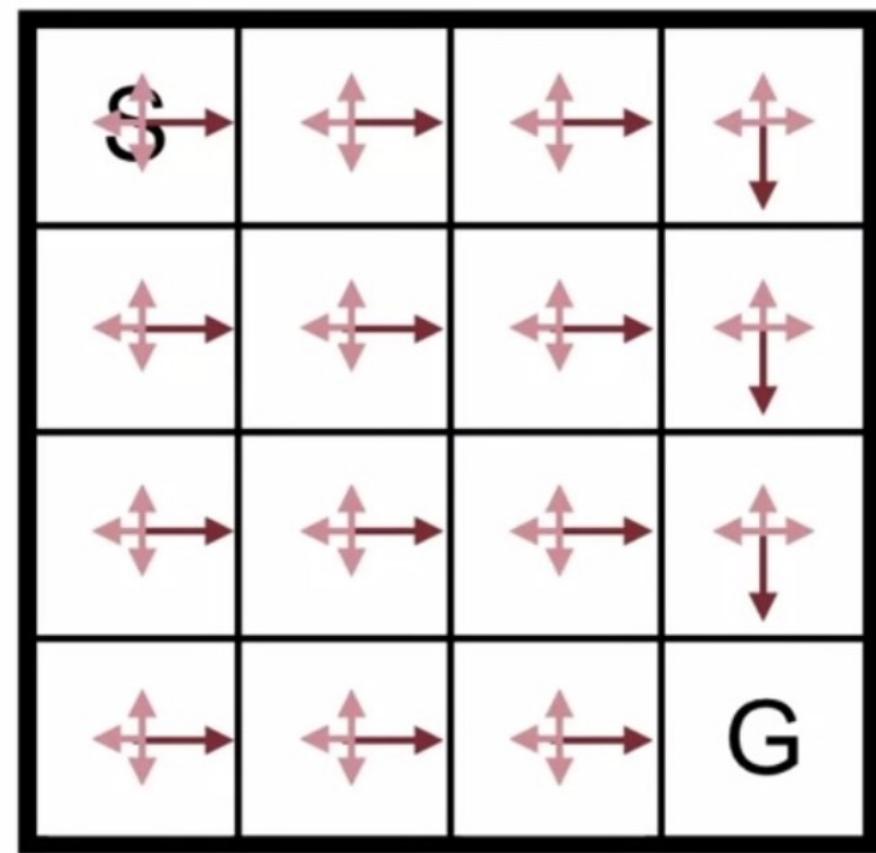
$$\pi(a|s) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|\mathcal{A}(s)|} & \text{if } a \in \mathcal{A}(s) \text{ is the greedy action} \\ \frac{\epsilon}{|\mathcal{A}(s)|} & \text{if } a \in \mathcal{A}(s) \text{ is a non-greedy action} \end{cases}$$

- On-policy first-visit MC control algorithm (described on the next slide) converges to the best  $\epsilon$ -soft policy.

# Illustration of $\epsilon$ -greedy policy



Optimal policy  $\pi^*$



$\epsilon$ -greedy policy

# On-policy first-visit MC control (for $\epsilon$ -soft policies)

- Initialize:
  - A small parameter  $\epsilon > 0$
  - $\pi$  an arbitrary  $\epsilon$ -soft policy such that  $\pi(a|s) \geq \frac{\epsilon}{|\mathcal{A}(s)|}$  for all states and actions
  - $Q(s, a) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$
  - $Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- Loop forever (for each episode):
  - Generate an episode following  $\pi: S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
  - $G \leftarrow 0$
  - Loop for each step of episode,  $t = T - 1, T - 2, \dots, 0$ :
    - $G \leftarrow \gamma G + R_{t+1}$
    - If the pair  $S_t, A_t$  does not appear in  $S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}$ :
      - Append  $G$  to  $Returns(S_t, A_t)$
      - $Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$
      - $A^* \leftarrow \underset{a' \in A}{\operatorname{argmax}} Q(S_t, a')$  (with ties broken arbitrarily)
      - For all  $a \in \mathcal{A}(S_t)$   
$$\pi(a|S_t) = \begin{cases} 1 - \epsilon + \epsilon / |\mathcal{A}(S_t)| & \text{if } a = A^* \\ \epsilon / |\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

# Improvement of $\epsilon$ -greedy policy

- For any  $\epsilon$ -soft policy  $\pi$ , any  $\epsilon$ -greedy policy with respect to  $q_\pi$  is guaranteed to be better than or equal to  $\pi$ .
- Let  $\pi'$  be the  $\epsilon$ -greedy policy

$$\begin{aligned} q_\pi(s, \pi'(s)) &= \sum_{a \in A} \pi'(a|s) q_\pi(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}(s)|} \sum_{a \in A} q_\pi(s, a) + (1 - \epsilon) \max_{a' \in A} q_\pi(s, a') \\ &\geq \frac{\epsilon}{|\mathcal{A}(s)|} \sum_{a \in A} q_\pi(s, a) + (1 - \epsilon) \sum_{a \in A} \frac{\pi(a|s) - \frac{\epsilon}{|\mathcal{A}(s)|}}{1 - \epsilon} q_\pi(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}(s)|} \sum_{a \in A} q_\pi(s, a) - \frac{\epsilon}{|\mathcal{A}(s)|} \sum_{a \in A} q_\pi(s, a) + \sum_{a \in A} \pi(a|s) q_\pi(s, a) \\ &= v_\pi(s) \end{aligned}$$

since  $\sum_{a \in A} \pi(a|s) - \frac{\epsilon}{|\mathcal{A}(s)|} = 1 - \epsilon$ ,  $\pi(a|s) - \frac{\epsilon}{|\mathcal{A}(s)|} \geq 0$  since  $\pi$  is an  $\epsilon$ -soft policy, and  $\max_{a' \in A} q_\pi(s, a')$  larger than any convex combination of  $q_\pi(s, a)$

# Off-Policy

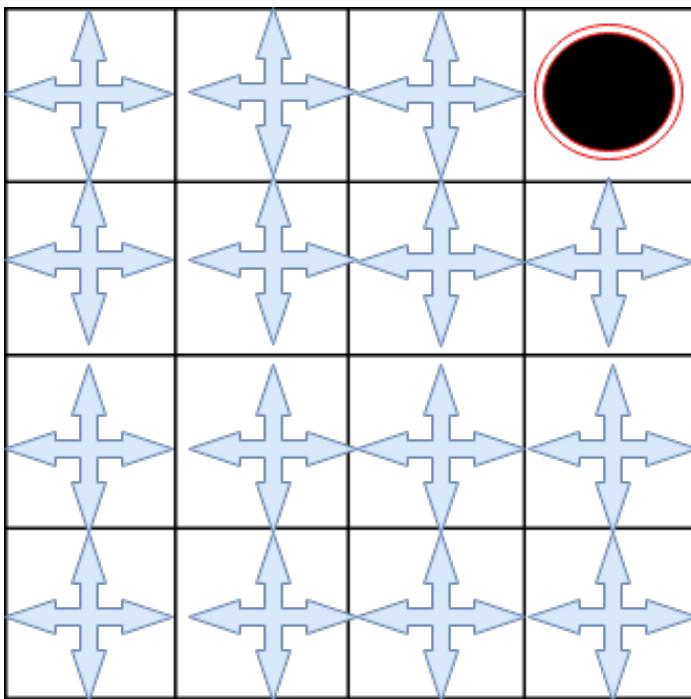
# Off-Policy Methods

- Learn the value of the target policy  $\pi$  from experience due to **behavior policy  $\mu$**
- For example,  $\pi$  is the greedy policy (and ultimately the optimal policy) while  $\mu$  is exploratory (e.g.,  $\epsilon$ -soft) policy
- In general, we only require **coverage**, i.e., that  $\mu$  generates behavior that covers or include  $\pi$

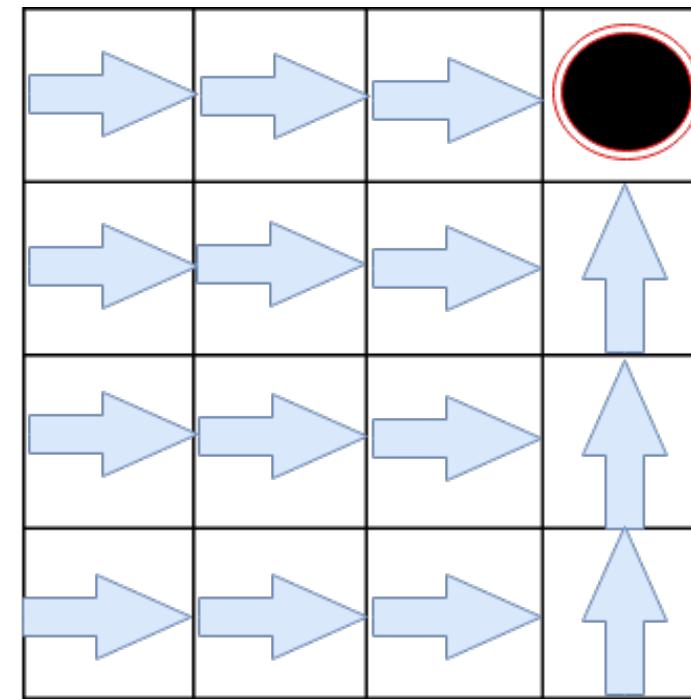
$$\mu(a|s) > 0 \text{ for every } s, a \text{ at which } \pi(a|s) > 0$$

- Can we average returns as before to obtain the value function of  $\pi$  ?
- Idea: Importance Sampling
  - Weight each return by the ratio of the probabilities of the trajectory under the two policies.

# Off-Policy Illustration



**Behavior Policy**



**Target Policy**

# Importance Sampling: Principle

- Estimate the expectation of a distribution  $\pi(X)$  from an other proposal distribution  $\mu(X)$

$$\mathbb{E}_{X \sim \pi}[f(X)] = \int f(X)\pi(X)dX = \int f(X) \frac{\pi(X)}{\mu(X)}\mu(X)dX = \mathbb{E}_{X \sim \mu} \left[ \frac{\pi(X)}{\mu(X)} f(X) \right]$$

- The two expectations are equivalent provided we use an appropriate weight
- Generally, we choose an easy-to-sample proposal distribution  $\mu(X)$
- In practice, we get

$$\mathbb{E}_{X \sim \pi}[f(X)] \approx \frac{1}{N} \sum_{i=1}^N \frac{\pi(x^{(i)})}{\mu(x^{(i)})} f(x^{(i)}) = \frac{1}{N} \sum_{i=1}^N w^{(i)} f(x^{(i)})$$

when we have  $N$  i.i.d. samples  $x^{(i)} \sim \mu(X)$  of  $X$  generated from  $\mu(X)$

- The quantities  $w^{(i)} = \frac{\pi(x^{(i)})}{\mu(x^{(i)})}$  are known as importance weights

# Importance Sampling Ratio for Off-Policy

- Probability of the trajectory, after state  $S_t$  until the terminal state  $S_T$ , under policy  $\pi$ :

$$\begin{aligned}\mathbb{P}(A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t, A_t \sim \pi, \dots, A_{T-1} \sim \pi) &= \pi(A_t | S_t) p(S_{t+1} | S_t, A_t) \pi(A_{t+1} | S_{t+1}) \cdots p(S_T | S_{T-1}, A_{T-1}) \\ &= \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)\end{aligned}$$

- **Importance sampling:** each return is weighted by the relative probability of the trajectory under the target and behavior policies

$$\rho_t^{T-1} = \frac{\prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)}{\prod_{k=t}^{T-1} \mu(A_k | S_k) p(S_{k+1} | S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k | S_k) p(S_{k+1} | S_k, A_k)}{\mu(A_k | S_k) p(S_{k+1} | S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k | S_k)}{\mu(A_k | S_k)}$$

# Importance Sample for Off-Policy

- Notations
  - To simplify notations, it is convenient to assume that **all the episodes are concatenated in a single sequence**
    - If the first episode ends in a terminal state at time  $t = 100$ , than the next episode begins at time  $t = 101$
  - Let  $\mathcal{T}(s)$  the set of time steps that were first visits to  $s$  within their episodes
  - Let  $T(t)$  denote the first time of termination following  $t$
  - Let  $G_t$  denote the return after  $t$  up through  $T(t)$  when we use the behavior policy  $\mu$  to choose the actions
- Remember that the **value function** (it is simpler to present that the state-action value function) is the expected return:
$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\mu[\rho_t^{T-1} G_t | S_t = s]$$
- Then
  - $\{G_t\}_{t \in \mathcal{T}(s)}$  are the returns that pertain to state  $s$
  - $\{\rho_t^{T(t)-1}\}_{t \in \mathcal{T}(s)}$  are the corresponding importance-sampling ratios
- Ordinary importance sampling estimate:

$$\hat{v}(s) = \frac{\sum_{t \in \mathcal{T}(s)} \rho_t^{T(t)-1} G_t}{|\mathcal{T}(s)|}$$

# Off-Policy Monte-Carlo in Practice

- Use returns generated from  $\mu$  to evaluate  $\pi$
- Multiply importance sampling corrections along whole episode

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)\pi(A_{t+1}|S_{t+1}) \cdots \pi(A_{T-1}|S_{T-1})}{\mu(A_t|S_t)\mu(A_{t+1}|S_{t+1}) \cdots \mu(A_{T-1}|S_{T-1})} G_t = \rho_t^{T-1} G_t$$

- Update value towards corrected return

$$\hat{v}(S_t) \leftarrow \hat{v}(S_t) + \alpha \left( G_t^{\pi/\mu} - \hat{v}(S_t) \right)$$

- Can not be used if  $\mu$  is zero when  $\pi$  is non-zero
- Importance sampling can dramatically increase variance

# Self-normalized importance sampling

- Often, the density we want to sample from is only known up to an unknown constant, i.e.,  $\mu(x) = \frac{\tilde{\mu}(x)}{Z_\mu}$  where  $\tilde{\mu}(x)$  is known, but  $Z_\mu$  is not. Of course  $Z_\mu$  is determined by the requirement that the integral of  $\mu(x)$  be 1, but we may not be able to compute the integral  $\int \tilde{\mu}(X)dX = Z_\mu$ .
- Let our proposal be of the form  $\mu(x) = \frac{\tilde{\mu}(x)}{Z_\mu}$

$$\mathbb{E}_{X \sim \pi}[f(X)] = \int f(X)\pi(X)dX = \int f(X)\frac{\pi(X)}{\mu(X)}\mu(X)dX = \frac{Z_\mu}{Z_\pi} \int f(X)\frac{\tilde{\pi}(X)}{\tilde{\mu}(X)}\mu(X)dX \approx \frac{1}{N} \frac{Z_\mu}{Z_\pi} \sum_{i=1}^N \frac{\tilde{\pi}(x^{(i)})}{\tilde{\mu}(x^{(i)})} f(x^{(i)}) = \frac{1}{N} \frac{Z_\mu}{Z_\pi} \sum_{i=1}^N w^{(i)} f(x^{(i)})$$

where  $\int \tilde{\pi}(X)dX = Z_\pi$  such that  $\tilde{\pi}(X)$  may also not normalized

- How to deal with the normalization constants? We note that

$$\frac{Z_\pi}{Z_\mu} = \frac{1}{Z_\mu} \int \tilde{\pi}(X)dX = \frac{1}{Z_\mu} \int \frac{\tilde{\pi}(X)}{\mu(X)}\mu(X)dX = \int \frac{\tilde{\pi}(X)}{\tilde{\mu}(X)}\mu(X)dX \approx \frac{1}{N} \sum_{i=1}^N \frac{\tilde{\pi}(x^{(i)})}{\tilde{\mu}(x^{(i)})} = \frac{1}{N} \sum_{i=1}^N w^{(i)}$$

- It follows that

$$\mathbb{E}_{X \sim \pi}[f(X)] \approx \sum_{i=1}^N \frac{w^{(i)}}{\sum_{i=1}^N w^{(i)}} f(x^{(i)}), \quad w^{(i)} = \frac{\tilde{\pi}(x^{(i)})}{\tilde{\mu}(x^{(i)})}, \quad x^{(i)} \sim \mu(X)$$

# Incremental implementation

- Let  $G_i(s)$  denote the return after  $t_i$  up through  $T(t_i)$  when we use the behavior policy  $\mu$  to choose the actions starting from state  $s$
- Let  $G_1(s), G_2(s), \dots, G_N(s)$  be the sequence of returns
- Remember that:

$$v_\pi(s) = \mathbb{E}_\mu[\rho_t^{T-1} G_t | S_t = s] \approx \hat{v}_{N+1}(s) = \sum_{i=1}^N \frac{w^{(i)}}{\sum_{i=1}^N w^{(i)}} G_i(s)$$
$$w^{(i)} = \rho_{t_i}^{T(t_i)-1}$$

- Incremental formula (used in practice for large  $N$ ):

$$\hat{v}_{n+1}(s) = \hat{v}_n(s) + \frac{w^{(n)}}{c^{(n)}} (G_n(s) - \hat{v}_n(s)), \quad n \geq 1$$
$$c^{(n+1)} = c^{(n)} + w^{(n+1)}$$
$$c^{(0)} = 0$$
$$\hat{v}_0(s) \in \mathbb{R}$$

# Which target function to estimate?

- The off-policy approach works both for estimating the value function and the state-action value function
- In practice the self-normalized importance sampling is preferred
  - It is biased
  - But its variance is smaller
- The next slide shows the full algorithm for estimating the state-action value function

# Off-policy MC prediction for estimating $Q \approx q_\pi$

- Input:  $\pi$  an arbitrary target policy
- Initialize:
  - $Q(s, a) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$
  - $C(s, a) \leftarrow 0$  for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$
- Loop forever (for each episode):
  - $\mu \leftarrow$  any policy with coverage  $\pi$
  - Generate an episode following  $\mu$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$
  - $G \leftarrow 0$
  - Loop for each step of episode,  $t = T - 1, T - 2, \dots, 0$ , while  $W \neq 0$ :
    - $G \leftarrow \gamma G + R_{t+1}$
    - $C(S_t, A_t) \leftarrow C(S_t, A_t) + W$
    - $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)}(G - Q(S_t, A_t))$
    - $W \leftarrow W \frac{\pi(A_t | S_t)}{\mu(A_t | S_t)}$

# Conclusion

# Monte Carlo Advantages

- MC has several advantages over dynamic programming:
  - Can learn directly from interaction with environment
  - No need for full models
- MC methods provide an alternate policy evaluation process
- One issue to watch for: maintaining sufficient exploration
- Looked at distinction between on-policy and off-policy methods